

Diana Marczuk

CSC 407 Systems 2 – Homework 1

4/18/2021

2. My code in main.c:

```
{
char line[MAX_LINE];
int entry;

//*****
//DIANA'S CODE
//*****
char *endcPtr;

do
{
printf("Please enter a %s number between %d and %d: \n", descriptionCPtr, low, high);
fgets(line, MAX_LINE, stdin);
entry = (int)strtol(line, &endcPtr, 10); //using base 10
}
while (entry < low || entry > high);
return (entry);
}

//*****
// END DIANA'S CODE
//*****
```

3. Answers

- a. 45.99 seconds
- b. 0.15 seconds

4. Answers:

- a. 23.57 seconds
- b. 0.07 seconds

5. **Option -O0 with a good algorithm will be faster.** The bad algorithm that can be referenced in this questions is generating a list instead of a binary tree. Whether doing the compiler optimization or not, a list will be over 300 times slower than attempting to generate a binary tree. Compiler optimization is not enough to account for the bad algorithm (should you choose to generate a linear list), as the optimization only makes the bad algorithm (and the good algorithm for that matter), twice as fast as when the programmer uses no compiler optimization, which is not enough to beat the time to generate a tree with the same data provided.

Should the algorithm be good, optimization can be added to make the process run even faster, however a programmer must watch out for unintentional changes the optimization makes to an algorithm.

6.

Question	Command	Result
(A) The string "%d: %d\n" is printed in printList()	objdump -s -j .rodata assign1-0	<pre>[dmarczuk@cdmlinux AssignmentOne]\$ objdump -s -j .rodata assign1-0 assign1-0: file format elf64-x86-64 Contents of section .rodata: 400db8 01000200 00000000 00000000 00000000 400dc8 506c6561 73652065 6e746572 20612025 Please enter a % 400dd8 73206e75 6d626572 20626574 7765656e s number between 400de8 20256420 616e6420 25643a20 0a000000 %d and %d: 400df8 74686520 6c6f7765 7374206e 756d6265 the lowest numbe 400e08 7220696e 20746865 2072616e 67650000 r in the range.. 400e18 74686520 68696768 65737420 6e756d62 the highest numb 400e28 65722069 6e207468 65207261 6e676500 er in the range. 400e38 74686520 6e756d62 6572206f 66206e75 the number of nu 400e48 6d626572 7320746f 20636f6e 73696465 mbers to conside 400e58 72000000 00000000 57686174 20776f75 r.....What wou 400e68 6c642079 6f75206c 696b6520 746f2064 ld you like to d 400e78 6f3f0a28 31292043 6f756e74 20776974 o2.(1) Count wit 400e88 68206120 6c697374 0a283229 20436f75 h a list.(2) Cou 400e98 6e742077 69746820 61207472 65650a28 nt with a tree.(400ea8 30292051 7569740a 596f7572 2063686f 0) Quit.Your cho 400eb8 69636520 0025643a 20256420 74696d65 ice %d: %d time 400ec8 2873290a 0025643a 20256420 74696d65 (s) %d: %d time 400ed8 2873290a 00 (s) ..</pre>
(B) The code for getNextNumber()	Objdump -d -j .text assign1-0	<pre>00000000004007ed <getNextNumber>: 4007ed: 55 push %rbp 4007ee: 48 89 e5 mov %rsp,%rbp 4007f1: e8 9a fe ff ff callq 400690 <__count@plt> 4007f6: e8 b5 fe ff ff callq 4006b0 <__rand@plt> 4007fb: 8b 0d 97 18 20 00 mov 0x201897(%rip),%ecx # 602098 <high> 400801: 8b 15 95 18 20 00 mov 0x201895(%rip),%edx # 60209c <low> 400807: 29 d1 sub %edx,%ecx 400809: 89 ca mov %ecx,%edx 40080b: 8d 4a 01 lea 0x1(%rdx),%ecx 40080e: 99 cltd 40080f: f7 f9 idiv %ecx 400811: 8b 05 85 18 20 00 mov 0x201885(%rip),%eax # 60209c <low> 400817: 01 d0 add %edx,%eax 400819: 5d pop %rbp 40081a: c3 retq</pre>
(C) The global variable high	Objdump -t -j .bss assign1-0	<pre>SYMBOL TABLE: 0000000000602080 l d .bss 0000000000000000 .bss 0000000000602090 l o .bss 0000000000000004 called.4239 0000000000602094 l o .bss 0000000000000001 completed.6355 0000000000602098 g o .bss 0000000000000004 high 00000000006020a0 g o .bss 0000000000000008 stdin@@GLIBC_2.2.5 00000000006020a0 g .bss 0000000000000000 _end 0000000000602074 g .bss 0000000000000000 __bss_start 000000000060209c g o .bss 0000000000000004 low</pre>
(D) treePtr in countWithTree()	objdump -s -j .data assign1-0	<pre>[dmarczuk@cdmlinux AssignmentOne]\$ objdump -s -j .data assign1-0 assign1-0: file format elf64-x86-64 Contents of section .data: 602070 00000000 TreePtr is a variable that has been declared but not defined yet, which is why it is zero. The compiler does not know where to actually put the variable yet because the end user did not put in a number since the program did not run yet, there it will put in zero for the time being.</pre>

7. **Optimization 1:** Variables are stored in rbp or ram in the not optimized code versus the optimized code (the second screenshot), which keeps many variables in their own registers.

```

40081a: c3          retq
NOT OPTIMIZED.....
00000000040081b <obtainNumberBetween>:
40081b: 55          push    %rbp
40081c: 48 89 e5    mov     %rsp,%rbp
40081f: 48 81 ec 30 01 00 00 sub     $0x130,%rsp
400826: e8 65 fe ff ff callq   400690 <mcount@plt>
40082b: 48 89 bd d8 fe ff ff mov     %rdi,-0x128(%rbp)
400832: 89 b5 d4 fe ff ff mov     %esi,-0x12c(%rbp)
400838: 89 95 d0 fe ff ff mov     %edx,-0x130(%rbp)
40083e: 8b d0 d0 fe ff ff mov     -0x130(%rbp),%ecx
400844: 8b 95 d4 fe ff ff mov     -0x12c(%rbp),%edx
40084a: 48 8b 85 d8 fe ff ff mov     -0x128(%rbp),%rax
400851: 48 89 c6    mov     %rax,%rsi
400854: bf c8 0d 40 00 mov     $0x400dc8,%edi
400859: b8 00 00 00 00 mov     $0x0,%eax
40085e: e8 bd fd ff ff callq   400620 <printf@plt>
400863: 48 8b 15 16 18 20 00 mov     0x201816(%rip),%rdx # 602080 <stdin@GLIBC_2.2.5>
40086a: 48 8d 85 f0 fe ff ff lea     -0x110(%rbp),%rax
400871: be 00 01 00 00 mov     $0x100,%esi
400876: 48 89 c7    mov     %rax,%rdi
400879: e8 d2 fd ff ff callq   400650 <fgets@plt>
40087e: 48 8d 8d e8 fe ff ff lea     -0x118(%rbp),%rcx
400885: 48 8d 85 f0 fe ff ff lea     -0x110(%rbp),%rax
40088c: ba 0a 00 00 00 mov     $0xa,%edx
400891: 48 89 ce    mov     %rcx,%rsi
400894: 48 89 c7    mov     %rax,%rdi
400897: e8 d4 fd ff ff callq   400670 <strtol@plt>
40089c: 89 45 fc    mov     %eax,-0x4(%rbp)
40089f: 8b 45 fc    mov     -0x4(%rbp),%eax
4008a2: 3b 85 d4 fe ff ff cmp     -0x12c(%rbp),%eax
4008a8: 7c 94      jl      40083e <obtainNumberBetween+0x23>
4008aa: 8b 45 fc    mov     -0x4(%rbp),%eax
4008ad: 3b 85 d0 fe ff ff cmp     -0x130(%rbp),%eax
4008b3: 7f 89      jg      40083e <obtainNumberBetween+0x23>
4008b5: 8b 45 fc    mov     -0x4(%rbp),%eax
4008b8: c9        leaveq
4008b9: c3          retq

```

```

0000000004008c0 <obtainNumberBetween>:
4008c0: 55          push    %rbp
4008c1: 48 89 e5    mov     %rsp,%rbp
4008c4: 41 55      push    %r13
4008c6: 41 54      push    %r12
4008c8: 53          push    %rbx
4008c9: 48 81 ec 18 01 00 00 sub     $0x118,%rsp
4008d0: e8 bb fd ff ff callq   400690 <mcount@plt>
4008d5: 49 89 fd    mov     %rdi,%r13
4008d8: 41 89 f4    mov     %esi,%r12d
4008db: 89 d3      mov     %edx,%ebx
4008dd: 0f 1f 00    nopl    (%rax)
4008e0: 89 d9      mov     %ebx,%ecx
4008e2: 44 89 e2    mov     %r12d,%edx
4008e5: 4c 89 ee    mov     %r13,%rsi
4008e8: bf 98 0d 40 00 mov     $0x400d98,%edi
4008ed: 31 c0      xor     %eax,%eax
4008ef: e8 2c fd ff ff callq   400620 <printf@plt>
4008f4: 48 8b 15 85 17 20 00 mov     0x201785(%rip),%rdx # 602080 <stdin@GLIBC_2.2.5>
4008fb: 48 8d bd e0 fe ff ff lea     -0x120(%rbp),%rdi
400902: be 00 01 00 00 mov     $0x100,%esi
400907: e8 44 fd ff ff callq   400650 <fgets@plt>
40090c: 48 8d b5 d8 fe ff ff lea     -0x128(%rbp),%rsi
400913: 48 8d bd e0 fe ff ff lea     -0x120(%rbp),%rdi
40091a: ba 0a 00 00 00 mov     $0xa,%edx
40091f: e8 4c fd ff ff callq   400670 <strtol@plt>
400924: 39 c3      cmp     %eax,%ebx
400926: 7c b8      jl      4008e0 <obtainNumberBetween+0x20>
400928: 41 39 c4    cmp     %eax,%r12d
40092b: 7f b3      jg      4008e0 <obtainNumberBetween+0x20>
40092d: 48 81 c4 18 01 00 00 add     $0x118,%rsp
400934: 5b          pop     %rbx
400935: 41 5c      pop     %r12
400937: 41 5d      pop     %r13
400939: 5d          pop     %rbp
40093a: c3          retq
40093b: 0f 1f 44 00 00 nopl    0x0(%rax,%rax,1)

```

Optimization 2: In tree.c, the generateTree() function, the non optimized version is allocating space on the stack for local variables instead of using registers r15 to r12 and rbx. The difference between these two sub commands between non optimized and optimized

version is 0x4 in the sub command. Sub command allocates less space on the stack for the optimization and instead uses registers to do operations, whereas the non optimized version uses more space on the stack and does not take advantage of registers.

No optimization-

```

00000000040096b <generateTree>:
40096b: 55                push    %rbp
40096c: 48 89 e5          mov     %rsp,%rbp
40096f: 48 83 ec 40       sub     $0x40,%rsp
400973: e8 18 fd ff ff    callq  400690 <mcount@plt>
400978: 89 7d cc          mov     %edi,-0x34(%rbp)
40097b: 48 c7 45 f8 00 00 movq    $0x0,-0x8(%rbp)
400982: 00
400983: c7 45 e4 00 00 00 movl    $0x0,-0x1c(%rbp)
40098a: e9 ec 00 00 00    jmpq   400a7b <generateTree+0x110>
40098f: b8 00 00 00 00    mov     $0x0,%eax
400994: e8 54 fe ff ff    callq  4007ed <getNextNumber>
400999: 89 45 e0          mov     %eax,-0x20(%rbp)
40099c: 48 c7 45 f0 00 00 movq    $0x0,-0x10(%rbp)
4009a3: 00
4009a4: 48 8b 45 f8       mov     -0x8(%rbp),%rax

```

With optimization-

```

000000000400940 <generateTree>:
400940: 55                push    %rbp
400941: 48 89 e5          mov     %rsp,%rbp
400944: 41 57             push    %r15
400946: 41 56             push    %r14
400948: 41 55             push    %r13
40094a: 41 54             push    %r12
40094c: 53               push    %rbx
40094d: 48 83 ec 08       sub     $0x8,%rsp
400951: e8 3a fd ff ff    callq  400690 <mcount@plt>
400956: 85 ff            test    %edi,%edi

```